

# FedSearch-NLP: Federated RAG QA System (Organization-Scale)

## 1. Project Overview

FedSearch-NLP is a large-scale, organization-level Retrieval-Augmented Question Answering (RAG) system built with Federated Learning (FL). Each department in an enterprise holds private documents (PDFs, SOPs, chats, reports). These departments train a shared global QA system collaboratively without sharing any raw text.

This system uses a centralized aggregator to merge encrypted updates, while all sensitive data stays within departments.

---

## 2. Core Architecture

### 2.1 Components per Client (Department)

- Local document preprocessing
- Local vector DB using FAISS
- Local embedding model (Sentence-Transformer / BERT-Large)
- Local adapter-based LLM for answer generation (Flan-T5-Large with LoRA)
- Local RAG pipeline
- Local training loop + DP-SGD

### 2.2 Central Server (Aggregator)

- Parameter aggregation (FedAvg / FedProx / SCAFFOLD)
- Handles secure aggregation
- Maintains global retriever + adapter weights

### 2.3 Privacy Protections

- Differential Privacy (DP-SGD)
  - Update encryption
  - No raw documents ever shared
- 

## 3. Model Choices (Large, Enterprise-Level)

### Retriever Model

- **all-mpnet-base-v2** or **bge-large-en** (enterprise-grade embedding models)
- Trained federatively across departments

## **Generator Model (LLM)**

- **Flan-T5-Large** (780M) with LoRA adapters
- Only LoRA layers are trained and federated
- Base model remains frozen for speed & stability

## **RAG Pipeline**

1. Query → client embedding
  2. FAISS search (local)
  3. Top-k context passed to generator
  4. Generator produces answer
- 

## **4. Workflow Summary**

1. Each department loads global model.
  2. Local RAG training on their private documents.
  3. LoRA & retriever updates computed.
  4. DP noise added.
  5. Secure aggregation merges updates globally.
  6. Updated global model redistributed.
- 

## **5. Federated Algorithms Used**

- **FedAvg** → baseline
  - **FedProx** → best for domain drift across departments
  - **Adapters-only FL** → reduces communication cost 95%
- 

## **6. Dataset Strategy**

### **Internal Documents (simulated)**

- HR policies
- Legal docs
- IT technical SOPs
- Slack/email chat logs
- Product documentation
- Research reports

### **Public Sources (for base fine-tuning)**

- Wikipedia dumps
- SQuAD / NQ

- MS MARCO
- 

## 7. Evaluation Metrics

### Retrieval

- Recall@5, Recall@10
- MRR

### QA

- F1
- EM (Exact Match)
- ROUGE-L

### Federated Metrics

- Communication cost per round
  - Accuracy vs number of clients
  - Privacy budget ( $\epsilon$ )
- 

## 8. Deployment Plan

### Phase 1: Local Simulator

- Multiple simulated clients on one machine
- Flower or FedML

### Phase 2: Cloud Deployment

- Docker containers (one per department)
  - Load balancer for server
- 

## 9. Code Components (that will be delivered next)

- **client.py** → preprocessing, FAISS index, local training
- **server.py** → FedAvg aggregation
- **model\_retriever.py** → embedding model
- **model\_generator.py** → T5 + LoRA
- **rag\_pipeline.py** → query → retrieve → generate
- **federated\_train.py** → full system runner
- **configs/** → hyperparameters
- **requirements.txt**

---

## 10. Final Output Deliverables

You will receive in next step:

- Complete working codebase (PyTorch + RAG + FL)
  - Ready-to-run simulation
  - Configurable number of departments/clients
  - Trained sample model
  - Startup instructions
- 

## 11. Notes

- Optimized for completion in 2 days
  - Minimal but powerful architecture
  - Uses enterprise-level models, not small ones
- 

## 12. Next Step

I will now generate the **full code** for: - server - client - local training - RAG modules - FL integration - FAISS indexing - LoRA adapters setup

Once code is integrated, I will also prepare a GitHub-ready README.